

510518079_Assignment3

September 21, 2021

1 Assignment 3 : Study on multiple datasets

1.1 Import Libraries

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import OneHotEncoder
from scipy.sparse import hstack
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix
from sklearn.preprocessing import StandardScaler,MinMaxScaler
```

1.2 Titanic Dataset

Number of features = 26 , excluding passenger ID and 2urvived, but the actual number of features are 7, as we can remove the zero columns as they are in here for One Hot Encoding. Number of data rows = 1309 but two rows contain NaN values, so after dropping them we get 1307 rows Numerical Attributes = Age of the passenger, Fare paid, 'Parch' or # of parents/children, 'sibsp' or # of siblings/spouses Categorical Attributes = 'Pclass' or class of passenger, sex of the passenger, 'Embarked' the port of Embarkation

1.2.1 Reading Dataset and Pre processing

```
[ ]: titanic = pd.read_csv('titanic.csv')
titanic.head()
```

```
[ ]: 
```

	Passengerid	Age	Fare	Sex	sibsp	zero	zero.1	zero.2	zero.3	\
0	1	22.0	7.2500	0	1	0	0	0	0	
1	2	38.0	71.2833	1	1	0	0	0	0	
2	3	26.0	7.9250	1	0	0	0	0	0	
3	4	35.0	53.1000	1	1	0	0	0	0	
4	5	35.0	8.0500	0	0	0	0	0	0	

	zero.4	...	zero.12	zero.13	zero.14	Pclass	zero.15	zero.16	Embarked	\
0	0	...	0	0	0	3	0	0	2.0	
1	0	...	0	0	0	1	0	0	0.0	
2	0	...	0	0	0	3	0	0	2.0	
3	0	...	0	0	0	1	0	0	2.0	
4	0	...	0	0	0	3	0	0	2.0	

	zero.17	zero.18	Survived
0	0	0	0
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	0

[5 rows x 28 columns]

```
[ ]: titanic.describe()
```

```
[ ]:
```

	Passengerid	Age	Fare	Sex	sibsp	\
count	1309.000000	1309.000000	1309.000000	1309.000000	1309.000000	
mean	655.000000	29.503186	33.281086	0.355997	0.498854	
std	378.020061	12.905241	51.741500	0.478997	1.041658	
min	1.000000	0.170000	0.000000	0.000000	0.000000	
25%	328.000000	22.000000	7.895800	0.000000	0.000000	
50%	655.000000	28.000000	14.454200	0.000000	0.000000	
75%	982.000000	35.000000	31.275000	1.000000	1.000000	
max	1309.000000	80.000000	512.329200	1.000000	8.000000	

	zero	zero.1	zero.2	zero.3	zero.4	...	zero.12	zero.13	zero.14	\
count	1309.0	1309.0	1309.0	1309.0	1309.0	...	1309.0	1309.0	1309.0	
mean	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
std	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
min	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
25%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
50%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
75%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
max	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	

	Pclass	zero.15	zero.16	Embarked	zero.17	zero.18	\
count	1309.000000	1309.0	1309.0	1307.000000	1309.0	1309.0	
mean	2.294882	0.0	0.0	1.492731	0.0	0.0	
std	0.837836	0.0	0.0	0.814626	0.0	0.0	
min	1.000000	0.0	0.0	0.000000	0.0	0.0	
25%	2.000000	0.0	0.0	1.000000	0.0	0.0	
50%	3.000000	0.0	0.0	2.000000	0.0	0.0	
75%	3.000000	0.0	0.0	2.000000	0.0	0.0	

```
max      3.000000      0.0      0.0      2.000000      0.0      0.0
```

```

      2survived
count  1309.000000
mean    0.261268
std     0.439494
min     0.000000
25%     0.000000
50%     0.000000
75%     1.000000
max     1.000000

```

```
[8 rows x 28 columns]
```

```
[ ]: titanic.columns
```

```
[ ]: Index(['Passengerid', 'Age', 'Fare', 'Sex', 'sibsp', 'zero', 'zero.1',
          'zero.2', 'zero.3', 'zero.4', 'zero.5', 'zero.6', 'Parch', 'zero.7',
          'zero.8', 'zero.9', 'zero.10', 'zero.11', 'zero.12', 'zero.13',
          'zero.14', 'Pclass', 'zero.15', 'zero.16', 'Embarked', 'zero.17',
          'zero.18', '2urvived'],
          dtype='object')
```

```
[ ]: titanic = titanic.dropna()
```

```
[ ]: X = titanic.copy()
for fn in X.columns:
    if len(X[fn].unique()) == 1:
        X = X.drop(columns = [fn])
Y = titanic['2urvived']
X = X.drop(columns = ['Passengerid', '2urvived'])
print(X.shape)
print(Y.shape)
X.head()
```

```
(1307, 7)
```

```
(1307,)
```

```
[ ]:
   Age    Fare  Sex  sibsp  Parch  Pclass  Embarked
0  22.0   7.2500   0     1     0       3         2.0
1  38.0  71.2833   1     1     0       1         0.0
2  26.0   7.9250   1     0     0       3         2.0
3  35.0  53.1000   1     1     0       1         2.0
4  35.0   8.0500   0     0     0       3         2.0
```

```
[ ]: scaler = MinMaxScaler()
```

```
X[['Age', 'Fare', 'Parch', 'sibsp']] = scaler.fit_transform(X[['Age', 'Fare', 'Parch', 'sibsp']])
X.head()
```

```
[ ]:      Age      Fare  Sex  sibsp  Parch  Pclass  Embarked
0  0.273456  0.014151   0  0.125   0.0      3        2.0
1  0.473882  0.139136   1  0.125   0.0      1        0.0
2  0.323563  0.015469   1  0.000   0.0      3        2.0
3  0.436302  0.103644   1  0.125   0.0      1        2.0
4  0.436302  0.015713   0  0.000   0.0      3        2.0
```

```
[ ]: pclass = X['Pclass'].values.reshape(-1, 1)
embarked = X['Embarked'].values.reshape(-1, 1)
sex = X['Sex'].values.reshape(-1, 1)

ohe = OneHotEncoder()

pclass_ohe = ohe.fit_transform(pclass)
embarked_ohe = ohe.fit_transform(embarked)
sex_ohe = ohe.fit_transform(sex)
```

```
[ ]: t_n = X[['Age', 'Fare', 'Parch', 'sibsp']]
#t_n = scaler.fit_transform
X = hstack((t_n, embarked_ohe, sex_ohe, pclass_ohe))
print(X.shape)
```

(1307, 12)

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify = Y,
↳test_size = 0.2, random_state = 5)
print(x_train.shape, x_test.shape)
print(y_train.shape, y_test.shape)
```

(1045, 12) (262, 12)

(1045,) (262,)

1.2.2

First, I normalised all the numerical features. Then, I converted categorical features by using One Hot Encoder Combined them together to get a matrix of dimension, (1307,12) Split the data into testing and training part.

1.2.3 Logistic Regression

```
[ ]: lr = LogisticRegression()
lr.fit(x_train, y_train)
y_pred = lr.predict(x_test)
lr_acc = accuracy_score(y_test, y_pred)
```

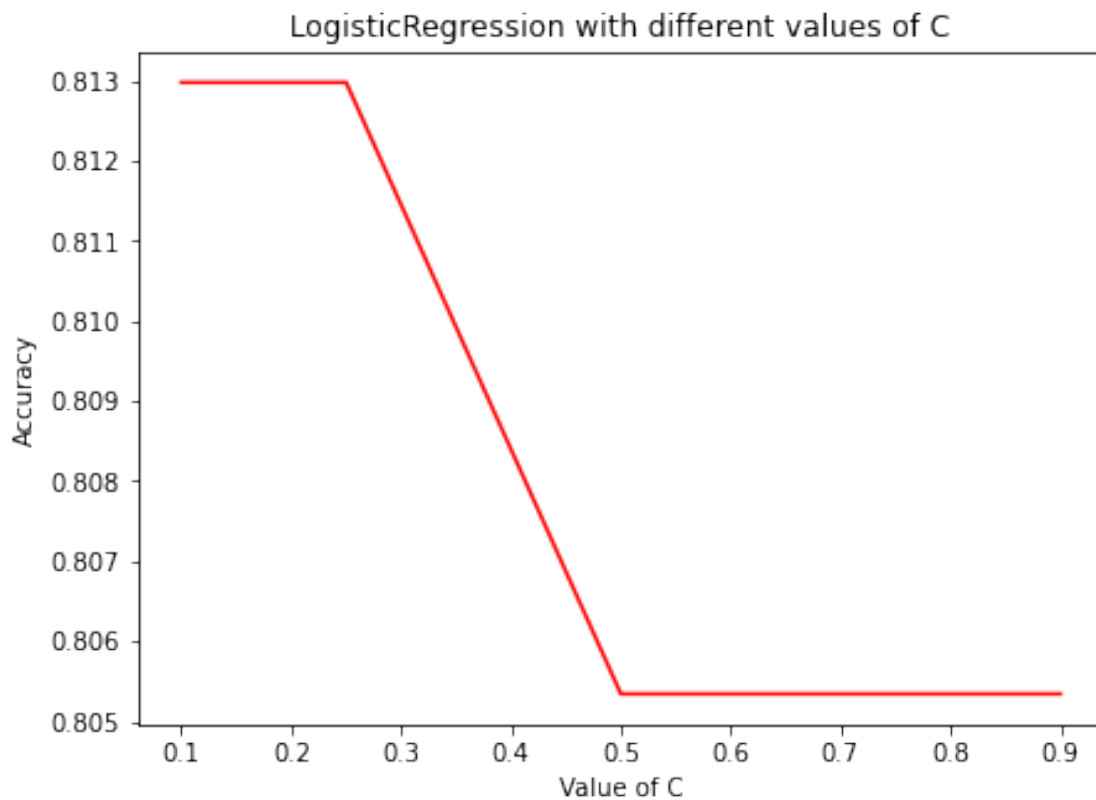
```
print(accuracy_score(y_test, y_pred))
```

0.8053435114503816

```
[ ]: cs = [0.1, 0.25, 0.5, 0.75, 0.9]
      coeffs = []
      accs = []

      for c in cs:
          lr = LogisticRegression(C = c)
          lr.fit(x_train, y_train)
          y_pred = lr.predict(x_test)
          coeffs.append(lr.coef_)
          lr_acc = accuracy_score(y_test, y_pred)
          accs.append(lr_acc)
```

```
[ ]: plt.figure(figsize = (7, 5))
      plt.title('LogisticRegression with different values of C')
      plt.ylabel('Accuracy')
      plt.xlabel('Value of C')
      plt.plot(cs, accs, color = 'red')
      plt.show()
```




After, performing the above experiment, I can clearly observe that best accuracy is around 81.3% at $C = 0.25$.

1.2.4 Decision Tree Classifier

```
[ ]: accs = []
for d in range(1,51,1):
    dt = DecisionTreeClassifier(max_depth = d)
    dt.fit(x_train, y_train)
    y_pred = dt.predict(x_test)
    dt_acc = accuracy_score(y_test, y_pred)
    accs.append(dt_acc)

[ ]: ds = [i for i in range(1,51,1)]
plt.figure(figsize = (10, 6))
plt.title('Accuracy of DecisionTreeClassifier with different values of max_depth')
plt.ylabel('Accuracy')
plt.xlabel('Max Depth')
plt.plot(ds, accs, color = 'blue')
plt.show()
```



I iterated from 1 to 50 and varied the max-depth in similar manner and from the following I can see that the best accuracy is at max depth 3, which is equal to ~81%. There is no trend that can

be observed on the basis of max-depth.

1.2.5 Conclusion

The Logistic Regression model performed better in the case of Titanic Data set for classification of survivor as compared to Decision Tree Classifier, although marginally.

1.3 Forest Cover Type Dataset

Part 1

There are 54 attributes excluding the cover_type , and in total of 581012 data rows. There are 14 numerical attributes and 40 categorical attributes, which are essentially soil types.

Part 2

There are total of 2 attributes excluding the cover type and 530895 data rows. There are only 2 numerical attributes and no categorical attributes.

1.3.1 Reading dataset and data pre processing

```
[ ]: forest = pd.read_csv("covtype.csv")
forest.head()
```

```
[ ]:
Elevation  Aspect  Slope  Horizontal_Distance_To_Hydrology  \
0      2596     51     3                      258
1      2590     56     2                      212
2      2804    139     9                      268
3      2785    155    18                      242
4      2595     45     2                      153

Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \
0                             0                             510
1                             -6                             390
2                             65                             3180
3                             118                            3090
4                             -1                             391

Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \
0             221             232             148
1             220             235             151
2             234             238             135
3             238             238             122
4             220             234             150

Horizontal_Distance_To_Fire_Points  ...  Soil_Type32  Soil_Type33  \
0             6279  ...             0             0
1             6225  ...             0             0
2             6121  ...             0             0
3             6211  ...             0             0
```

```

4                                6172 ...                                0                                0

Soil_Type34 Soil_Type35 Soil_Type36 Soil_Type37 Soil_Type38 \
0           0           0           0           0           0
1           0           0           0           0           0
2           0           0           0           0           0
3           0           0           0           0           0
4           0           0           0           0           0

```

```

Soil_Type39 Soil_Type40 Cover_Type
0           0           0           5
1           0           0           5
2           0           0           2
3           0           0           2
4           0           0           5

```

[5 rows x 55 columns]

```
[ ]: forest.describe()
```

```

[ ]:      Elevation      Aspect      Slope \
count  581012.000000  581012.000000  581012.000000
mean    2959.365301    155.656807    14.103704
std      279.984734    111.913721     7.488242
min     1859.000000     0.000000     0.000000
25%     2809.000000     58.000000     9.000000
50%     2996.000000    127.000000    13.000000
75%     3163.000000    260.000000    18.000000
max     3858.000000    360.000000    66.000000

```

```

      Horizontal_Distance_To_Hydrology  Vertical_Distance_To_Hydrology \
count                                581012.000000                    581012.000000
mean                                  269.428217                      46.418855
std                                   212.549356                      58.295232
min                                   0.000000                      -173.000000
25%                                   108.000000                       7.000000
50%                                   218.000000                      30.000000
75%                                   384.000000                      69.000000
max                                   1397.000000                     601.000000

```

```

      Horizontal_Distance_To_Roadways  Hillshade_9am  Hillshade_Noon \
count                                581012.000000    581012.000000    581012.000000
mean                                  2350.146611      212.146049      223.318716
std                                   1559.254870      26.769889      19.768697
min                                   0.000000         0.000000         0.000000
25%                                   1106.000000      198.000000      213.000000
50%                                   1997.000000      218.000000      226.000000

```


75%	3328.000000	231.000000	237.000000
max	7117.000000	254.000000	254.000000

	Hillshade_3pm	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	\
count	581012.000000	581012.000000	...	581012.000000	
mean	142.528263	1980.291226	...	0.090392	
std	38.274529	1324.195210	...	0.286743	
min	0.000000	0.000000	...	0.000000	
25%	119.000000	1024.000000	...	0.000000	
50%	143.000000	1710.000000	...	0.000000	
75%	168.000000	2550.000000	...	0.000000	
max	254.000000	7173.000000	...	1.000000	

	Soil_Type33	Soil_Type34	Soil_Type35	Soil_Type36	\
count	581012.000000	581012.000000	581012.000000	581012.000000	
mean	0.077716	0.002773	0.003255	0.000205	
std	0.267725	0.052584	0.056957	0.014310	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	Soil_Type37	Soil_Type38	Soil_Type39	Soil_Type40	\
count	581012.000000	581012.000000	581012.000000	581012.000000	
mean	0.000513	0.026803	0.023762	0.015060	
std	0.022641	0.161508	0.152307	0.121791	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	Cover_Type
count	581012.000000
mean	2.051471
std	1.396504
min	1.000000
25%	1.000000
50%	2.000000
75%	2.000000
max	7.000000

[8 rows x 55 columns]

```
[ ]: forest.columns
```

```
[ ]: Index(['Elevation', 'Aspect', 'Slope', 'Horizontal_Distance_To_Hydrology',
          'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
          'Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm',
          'Horizontal_Distance_To_Fire_Points', 'Wilderness_Area1',
          'Wilderness_Area2', 'Wilderness_Area3', 'Wilderness_Area4',
          'Soil_Type1', 'Soil_Type2', 'Soil_Type3', 'Soil_Type4', 'Soil_Type5',
          'Soil_Type6', 'Soil_Type7', 'Soil_Type8', 'Soil_Type9', 'Soil_Type10',
          'Soil_Type11', 'Soil_Type12', 'Soil_Type13', 'Soil_Type14',
          'Soil_Type15', 'Soil_Type16', 'Soil_Type17', 'Soil_Type18',
          'Soil_Type19', 'Soil_Type20', 'Soil_Type21', 'Soil_Type22',
          'Soil_Type23', 'Soil_Type24', 'Soil_Type25', 'Soil_Type26',
          'Soil_Type27', 'Soil_Type28', 'Soil_Type29', 'Soil_Type30',
          'Soil_Type31', 'Soil_Type32', 'Soil_Type33', 'Soil_Type34',
          'Soil_Type35', 'Soil_Type36', 'Soil_Type37', 'Soil_Type38',
          'Soil_Type39', 'Soil_Type40', 'Cover_Type'],
          dtype='object')
```

```
[ ]: forest['Cover_Type'].value_counts()
```

```
[ ]: 2    283301
      1    211840
      3    35754
      7    20510
      6    17367
      5     9493
      4     2747
      Name: Cover_Type, dtype: int64
```

```
[ ]: forest['Soil_Type32'].value_counts()
```

```
[ ]: 0    528493
      1    52519
      Name: Soil_Type32, dtype: int64
```

```
[ ]: X1 = forest.copy()
      X1[['Elevation', 'Aspect', 'Slope', 'Horizontal_Distance_To_Hydrology',
          'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
          'Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm',
          'Horizontal_Distance_To_Fire_Points', 'Wilderness_Area1',
          'Wilderness_Area2', 'Wilderness_Area3', 'Wilderness_Area4']] = scaler.
      →fit_transform(X1[['Elevation', 'Aspect', 'Slope',
      →'Horizontal_Distance_To_Hydrology',
          'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
          'Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm',
          'Horizontal_Distance_To_Fire_Points', 'Wilderness_Area1',
          'Wilderness_Area2', 'Wilderness_Area3', 'Wilderness_Area4']])
      X1 = X1.drop(columns = ['Cover_Type'])
```

```
Y1 = forest['Cover_Type']
print(Y1.shape)
X1.head()
```

(581012,)

```
[ ]:  Elevation    Aspect    Slope  Horizontal_Distance_To_Hydrology  \
0    0.368684  0.141667  0.045455                                0.184681
1    0.365683  0.155556  0.030303                                0.151754
2    0.472736  0.386111  0.136364                                0.191840
3    0.463232  0.430556  0.272727                                0.173228
4    0.368184  0.125000  0.030303                                0.109520

      Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \
0                                0.223514                        0.071659
1                                0.215762                        0.054798
2                                0.307494                        0.446817
3                                0.375969                        0.434172
4                                0.222222                        0.054939

      Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \
0          0.870079          0.913386          0.582677
1          0.866142          0.925197          0.594488
2          0.921260          0.937008          0.531496
3          0.937008          0.937008          0.480315
4          0.866142          0.921260          0.590551

      Horizontal_Distance_To_Fire_Points  ...  Soil_Type31  Soil_Type32  \
0                                0.875366  ...           0           0
1                                0.867838  ...           0           0
2                                0.853339  ...           0           0
3                                0.865886  ...           0           0
4                                0.860449  ...           0           0

      Soil_Type33  Soil_Type34  Soil_Type35  Soil_Type36  Soil_Type37  \
0              0           0           0           0           0
1              0           0           0           0           0
2              0           0           0           0           0
3              0           0           0           0           0
4              0           0           0           0           0

      Soil_Type38  Soil_Type39  Soil_Type40
0              0           0           0
1              0           0           0
2              0           0           0
3              0           0           0
4              0           0           0
```

[5 rows x 54 columns]

```
[ ]: forestx = forest[(forest['Cover_Type'] == 1) | (forest['Cover_Type'] == 2) |  
    ↳(forest['Cover_Type'] == 3)]  
print(forestx.shape)  
forestx.head()
```

(530895, 55)

```
[ ]:      Elevation  Aspect  Slope  Horizontal_Distance_To_Hydrology  \  
2          2804      139      9                                268  
3          2785      155      18                                242  
5          2579      132      6                                 300  
11         2886      151      11                                371  
12         2742      134      22                                150  
  
      Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \  
2                                65                                3180  
3                               118                                3090  
5                               -15                                 67  
11                              26                               5253  
12                              69                               3215  
  
      Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \  
2              234              238              135  
3              238              238              122  
5              230              237              140  
11             234              240              136  
12             248              224              92  
  
      Horizontal_Distance_To_Fire_Points  ...  Soil_Type32  Soil_Type33  \  
2                                6121  ...           0           0  
3                                6211  ...           0           0  
5                                6031  ...           0           0  
11                               4051  ...           0           0  
12                               6091  ...           0           0  
  
      Soil_Type34  Soil_Type35  Soil_Type36  Soil_Type37  Soil_Type38  \  
2              0              0              0              0              0  
3              0              0              0              0              0  
5              0              0              0              0              0  
11             0              0              0              0              0  
12             0              0              0              0              0  
  
      Soil_Type39  Soil_Type40  Cover_Type  
2              0              0              2
```

3	0	0	2
5	0	0	2
11	0	0	2
12	0	0	2

[5 rows x 55 columns]

```
[ ]: forestx = forestx[['Slope', 'Elevation', 'Cover_Type']]
print(forestx.shape)
```

(530895, 3)

```
[ ]: X2 = forestx.copy()
X2[['Slope', 'Elevation']] = scaler.fit_transform(X2[['Slope', 'Elevation']])
X2 = X2.drop(columns = ['Cover_Type'])
Y2 = forestx['Cover_Type']

print(Y2.shape)
X2.head()
```

(530895,)

```
[ ]:      Slope  Elevation
2    0.136364  0.517241
3    0.272727  0.506842
5    0.090909  0.394089
11   0.166667  0.562124
12   0.333333  0.483306
```

```
[ ]: x_train1, x_test1, y_train1, y_test1 = train_test_split(X1, Y1, stratify = Y1,
↳test_size = 0.98, random_state = 42)
print(x_train1.shape, x_test1.shape)
print(y_train1.shape, y_test1.shape)
```

(11620, 54) (569392, 54)

(11620,) (569392,)

```
[ ]: x_train2, x_test2, y_train2, y_test2 = train_test_split(X2, Y2,
↳stratify=Y2, test_size = 0.2)
print(x_train2.shape, x_test2.shape)
print(y_train2.shape, y_test2.shape)
```

(424716, 2) (106179, 2)

(424716,) (106179,)

Part 1

I normalised my numerical features, then kept 2% of the dataset for training and 98% for testing.

Part 2

I just normalised my data and then split the data into training and testing sets.

1.3.2 SVM

```
[ ]: svc = SVC(verbose=True)
      svc.fit(x_train1, y_train1)
      y_pred1 = svc.predict(x_test1)
```

```
[ ]: print(accuracy_score(y_test1, y_pred1))
```


0.725127855677635

```
[ ]: cm = confusion_matrix(y_test1, y_pred1, labels=svc.classes_)
      print(cm)
```

```
[[143899  58347   180     0     0     0  5177]
 [ 44981 225180  6796     9     0    36   633]
 [     0   3841 30709   276     0   213     0]
 [     0     0  2064   425     0   203     0]
 [   167   8310   826     0     0     0     0]
 [     0   6046 10391    67     0   516     0]
 [   7733   137    77     0     0     0 12153]]
```

```
[ ]: ax = sns.heatmap(cm, annot=True, fmt='g')
      ax.set_title('Confusion Matrix')
      ax.set_xlabel('Predicted Labels')
      ax.set_ylabel('Actual Labels')
      ax.xaxis.set_ticklabels(svc.classes_)
      ax.yaxis.set_ticklabels(svc.classes_)
```

```
[ ]: [Text(0, 0.5, '1'),
      Text(0, 1.5, '2'),
      Text(0, 2.5, '3'),
      Text(0, 3.5, '4'),
      Text(0, 4.5, '5'),
      Text(0, 5.5, '6'),
      Text(0, 6.5, '7')]
```



Above, I have plotted the confusion matrix from where I can observe number of points for each combination of original and predicted class. We can see that our model is not able to classify a single point of label 5 and the classification from label 4 onwards is pretty poor. The accuracy of the model is 72.5%.

1.3.3 Logistic Regression

```
[ ]: logreg = LogisticRegression()
logreg.fit(x_train2, y_train2)
```

```
[ ]: LogisticRegression()
```

```
[ ]: x_min, x_max = min(X2.values[:, 0]) - 0.1, max(X2.values[:, 0]) + 0.1
y_min, y_max = min(X2.values[:, 1]) - 0.1, max(X2.values[:, 1]) + 0.1
h = .01 # step size in the mesh
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
Z = logreg.predict(np.c_[xx.ravel(), yy.ravel()])

# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.figure(1, figsize=(10, 8))
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)

# Plot also the training points
```


```
plt.title('1. Logistic Regression decision boundaries on with Training set')
plt.scatter(x_train2.values[:, 0], x_train2.values[:, 1], c=y_train2,
            edgecolors='k', cmap=plt.cm.Paired)
plt.xlabel('Slope')
plt.ylabel('Elevation')

plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())

plt.show()
```

/tmp/ipykernel_4194/3129041014.py:10: MatplotlibDeprecationWarning: shading='flat' when X and Y have the same dimensions as C is deprecated since 3.3. Either specify the corners of the quadrilaterals with X and Y, or pass shading='auto', 'nearest' or 'gouraud', or set rcParams['pcolor.shading']. This will become an error two minor releases later.

```
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)
```




```
[ ]: plt.figure(1, figsize=(10, 8))
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)

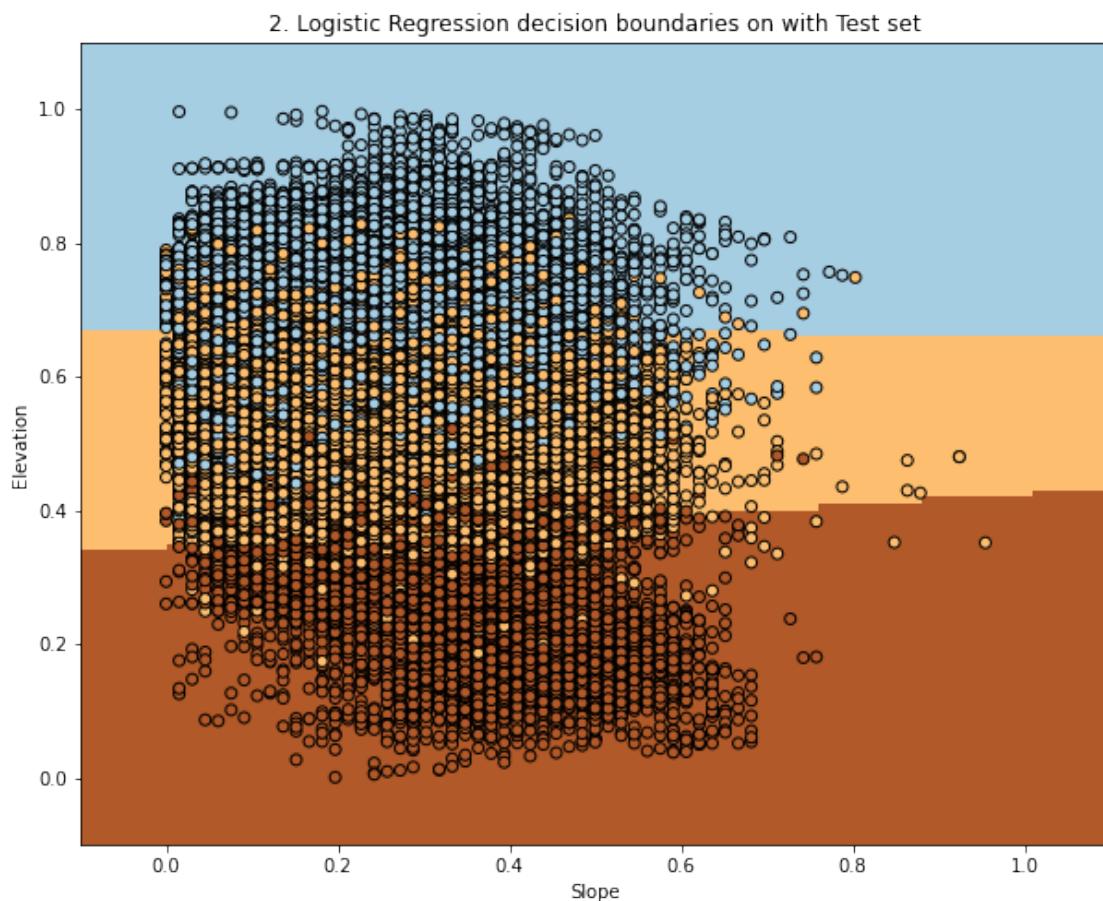
# Plot also the training points
plt.title('2. Logistic Regression decision boundaries on with Test set')
plt.scatter(x_test2.values[:, 0], x_test2.values[:, 1], c=y_test2,
            edgecolors='k', cmap=plt.cm.Paired)
plt.xlabel('Slope')
plt.ylabel('Elevation')

plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())

plt.show()
```

/tmp/ipykernel_4194/4182778859.py:2: MatplotlibDeprecationWarning:
shading='flat' when X and Y have the same dimensions as C is deprecated since
3.3. Either specify the corners of the quadrilaterals with X and Y, or pass
shading='auto', 'nearest' or 'gouraud', or set rcParams['pcolor.shading']. This
will become an error two minor releases later.

```
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)
```



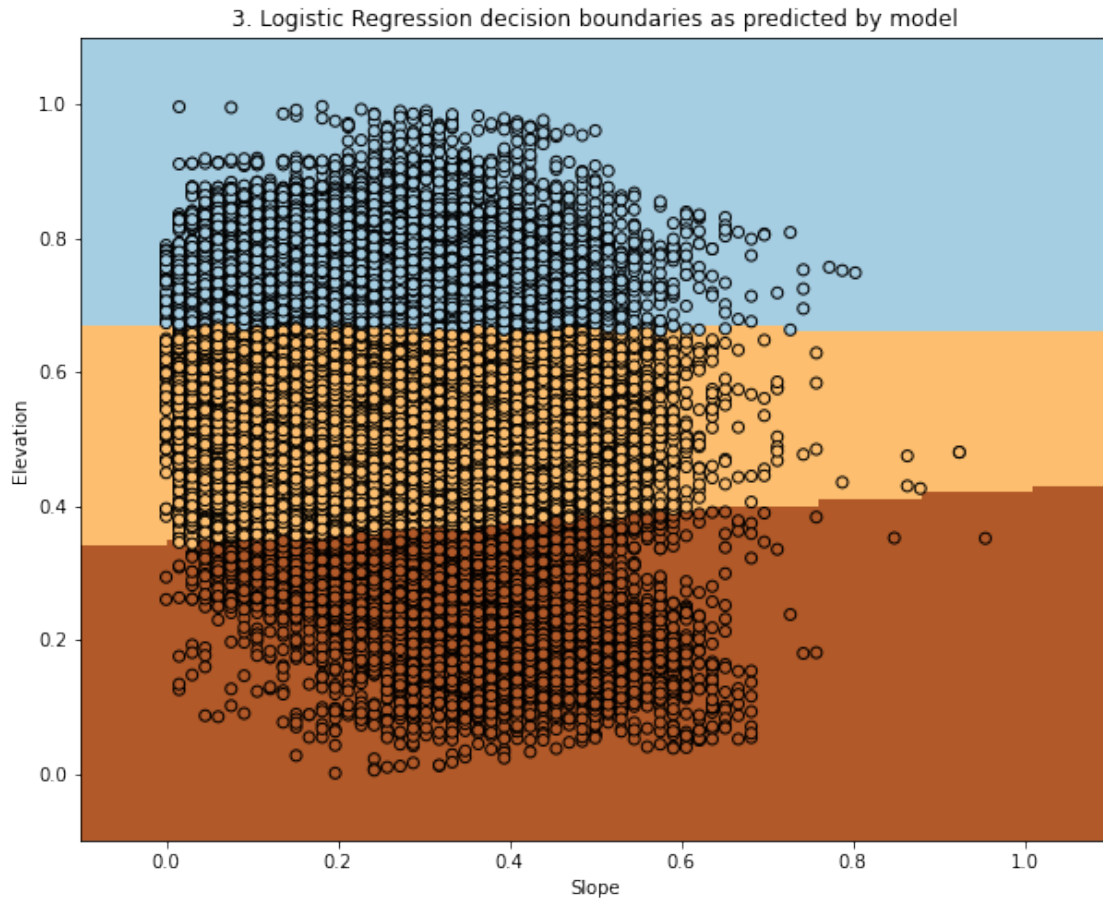
```
[ ]: y_pred2 = logreg.predict(x_test2)

plt.figure(1, figsize=(10, 8))
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)

# Plot also the training points
plt.title('3. Logistic Regression decision boundaries as predicted by model')
plt.scatter(x_test2.values[:, 0], x_test2.values[:, 1], c=y_pred2,
            ↪edgecolors='k', cmap=plt.cm.Paired)
plt.xlabel('Slope')
plt.ylabel('Elevation')
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())

plt.show()
```

/tmp/ipykernel_4194/3546256627.py:4: MatplotlibDeprecationWarning:
shading='flat' when X and Y have the same dimensions as C is deprecated since
3.3. Either specify the corners of the quadrilaterals with X and Y, or pass
shading='auto', 'nearest' or 'gouraud', or set rcParams['pcolor.shading']. This
will become an error two minor releases later.
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)



```
[ ]: print(accuracy_score(y_train2,logreg.predict(x_train2)))  
      print(accuracy_score(y_test2,y_pred2))
```

```
0.7333041373529605  
0.7345143578297026
```

From the plot 1, we can visualize the boundaries on the training data set. Plot 2 and Plot 3 gives us a Reality vs Expectation check for our model, where we can see what original distribution was and how are model perceived it. Also, we get a 73.3% accuracy on training dataset and 73.4% accuracy on testing data set.