

510518079_Assignment4

October 26, 2021

1 Assignment 4 : SVM

1.1 Import Libraries

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
from scipy.sparse import hstack
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.svm import SVC
from sklearn.multiclass import OneVsRestClassifier
from sklearn.multiclass import OneVsOneClassifier
from sklearn.preprocessing import MinMaxScaler
```

1.2 Pre Processing on Titanic Data Set

```
[ ]: titanic = pd.read_csv('titanic.csv')
titanic.head()
```

```
[ ]: 
```

	Passengerid	Age	Fare	Sex	sibsp	zero	zero.1	zero.2	zero.3	\
0	1	22.0	7.2500	0	1	0	0	0	0	
1	2	38.0	71.2833	1	1	0	0	0	0	
2	3	26.0	7.9250	1	0	0	0	0	0	
3	4	35.0	53.1000	1	1	0	0	0	0	
4	5	35.0	8.0500	0	0	0	0	0	0	

	zero.4	...	zero.12	zero.13	zero.14	Pclass	zero.15	zero.16	Embarked	\
0	0	...	0	0	0	3	0	0	2.0	
1	0	...	0	0	0	1	0	0	0.0	
2	0	...	0	0	0	3	0	0	2.0	
3	0	...	0	0	0	1	0	0	2.0	
4	0	...	0	0	0	3	0	0	2.0	

	zero.17	zero.18	Survived
--	---------	---------	----------

0	0	0	0
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	0

[5 rows x 28 columns]

```
[ ]: titanic.describe()
```

```
[ ]:
```

	Passengerid	Age	Fare	Sex	sibsp \
count	1309.000000	1309.000000	1309.000000	1309.000000	1309.000000
mean	655.000000	29.503186	33.281086	0.355997	0.498854
std	378.020061	12.905241	51.741500	0.478997	1.041658
min	1.000000	0.170000	0.000000	0.000000	0.000000
25%	328.000000	22.000000	7.895800	0.000000	0.000000
50%	655.000000	28.000000	14.454200	0.000000	0.000000
75%	982.000000	35.000000	31.275000	1.000000	1.000000
max	1309.000000	80.000000	512.329200	1.000000	8.000000

	zero	zero.1	zero.2	zero.3	zero.4	...	zero.12	zero.13	zero.14 \
count	1309.0	1309.0	1309.0	1309.0	1309.0	...	1309.0	1309.0	1309.0
mean	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
std	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
min	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
25%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
50%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
75%	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
max	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0

	Pclass	zero.15	zero.16	Embarked	zero.17	zero.18 \
count	1309.000000	1309.0	1309.0	1307.000000	1309.0	1309.0
mean	2.294882	0.0	0.0	1.492731	0.0	0.0
std	0.837836	0.0	0.0	0.814626	0.0	0.0
min	1.000000	0.0	0.0	0.000000	0.0	0.0
25%	2.000000	0.0	0.0	1.000000	0.0	0.0
50%	3.000000	0.0	0.0	2.000000	0.0	0.0
75%	3.000000	0.0	0.0	2.000000	0.0	0.0
max	3.000000	0.0	0.0	2.000000	0.0	0.0

	Survived
count	1309.000000
mean	0.261268
std	0.439494
min	0.000000
25%	0.000000
50%	0.000000

```
75%      1.000000
max      1.000000
```

```
[8 rows x 28 columns]
```

```
[ ]: titanic = titanic.dropna()
```

```
[ ]: X = titanic.copy()
for fn in X.columns:
    if len(X[fn].unique()) == 1:
        X = X.drop(columns = [fn])
Y = titanic['Survived']
X = X.drop(columns = ['PassengerId', 'Survived'])
print(X.shape)
print(Y.shape)
X.head()
```

```
(1307, 7)
```

```
(1307,)
```

```
[ ]:      Age      Fare  Sex  sibsp  Parch  Pclass  Embarked
0  22.0    7.2500    0     1      0      3         2.0
1  38.0   71.2833    1     1      0      1         0.0
2  26.0    7.9250    1     0      0      3         2.0
3  35.0   53.1000    1     1      0      1         2.0
4  35.0    8.0500    0     0      0      3         2.0
```

```
[ ]: scaler = MinMaxScaler()
X[['Age', 'Fare', 'Parch', 'sibsp']] = scaler.fit_transform(X[['Age', 'Fare', 'Parch', 'sibsp']])
X.head()
```

```
[ ]:      Age      Fare  Sex  sibsp  Parch  Pclass  Embarked
0  0.273456  0.014151    0  0.125    0.0      3         2.0
1  0.473882  0.139136    1  0.125    0.0      1         0.0
2  0.323563  0.015469    1  0.000    0.0      3         2.0
3  0.436302  0.103644    1  0.125    0.0      1         2.0
4  0.436302  0.015713    0  0.000    0.0      3         2.0
```

```
[ ]: pclass = X['Pclass'].values.reshape(-1, 1)
embarked = X['Embarked'].values.reshape(-1, 1)
sex = X['Sex'].values.reshape(-1, 1)

ohe = OneHotEncoder()

pclass_ohe = ohe.fit_transform(pclass)
embarked_ohe = ohe.fit_transform(embarked)
```

```
sex_ohe = ohe.fit_transform(sex)
```

```
[ ]: t_n = X[['Age', 'Fare', 'Parch', 'sibsp']]
      X = hstack((t_n, embarked_ohe, sex_ohe, pclass_ohe))
      print(X.shape)
```

```
(1307, 12)
```

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify = Y,
      ↳ test_size = 0.2, random_state = 10)
      print(x_train.shape, x_test.shape)
      print(y_train.shape, y_test.shape)
```

```
(1045, 12) (262, 12)
```

```
(1045,) (262,)
```

1.3 Task 1

```
[ ]: kernel, degree, n_0, n_1, acc = [], [], [], [], []
```

```
[ ]: model = SVC(kernel = 'linear')
      model.fit(x_train, y_train)
      y_pred = model.predict(x_test)
      print(accuracy_score(y_test, y_pred))
      print(model.n_support_)
      kernel.append('Linear')
      degree.append('NA')
      n_0.append(model.n_support_[0])
      n_1.append(model.n_support_[1])
      acc.append(accuracy_score(y_test, y_pred))
```

```
0.7938931297709924
```

```
[271 260]
```

```
[ ]: model = SVC(kernel = 'poly', degree = 2)
      model.fit(x_train, y_train)
      y_pred = model.predict(x_test)
      print(accuracy_score(y_test, y_pred))
      print(model.n_support_)
      kernel.append('Polynomial')
      degree.append('2')
      n_0.append(model.n_support_[0])
      n_1.append(model.n_support_[1])
      acc.append(accuracy_score(y_test, y_pred))
```

```
0.7938931297709924
```

```
[245 228]
```

```
[ ]: model = SVC(kernel = 'poly', degree = 3)
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(accuracy_score(y_test, y_pred))
print(model.n_support_)
kernel.append('Polynomial')
degree.append('3')
n_0.append(model.n_support_[0])
n_1.append(model.n_support_[1])
acc.append(accuracy_score(y_test, y_pred))
```

0.7824427480916031
[246 223]

```
[ ]: model = SVC(kernel = 'poly', degree = 5)
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(accuracy_score(y_test, y_pred))
print(model.n_support_)
kernel.append('Polynomial')
degree.append('5')
n_0.append(model.n_support_[0])
n_1.append(model.n_support_[1])
acc.append(accuracy_score(y_test, y_pred))
```

0.7786259541984732
[256 222]

```
[ ]: model = SVC(kernel = 'rbf')
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(accuracy_score(y_test, y_pred))
print(model.n_support_)
kernel.append('RBF')
degree.append('NA')
n_0.append(model.n_support_[0])
n_1.append(model.n_support_[1])
acc.append(accuracy_score(y_test, y_pred))
```

0.7748091603053435
[249 227]

```
[ ]: model = SVC(kernel = 'sigmoid')
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(accuracy_score(y_test, y_pred))
print(model.n_support_)
kernel.append('Sigmoid')
```

```
degree.append('NA')
n_0.append(model.n_support_[0])
n_1.append(model.n_support_[1])
acc.append(accuracy_score(y_test, y_pred))
```

```
0.6564885496183206
[205 204]
```

```
[ ]: comp = pd.DataFrame(zip(kernel, degree, n_0, n_1, acc), columns = ['Kernel', 'Degree', 'Class 0 SVs', 'Class 1 SVs', 'Accuracy'])
comp
```

```
[ ]:
```

	Kernel	Degree	Class 0 SVs	Class 1 SVs	Accuracy
0	Linear	NA	271	260	0.793893
1	Polynomial	2	245	228	0.793893
2	Polynomial	3	246	223	0.782443
3	Polynomial	5	256	222	0.778626
4	RBF	NA	249	227	0.774809
5	Sigmoid	NA	205	204	0.656489

1.4 Task 2

```
[ ]: titanic_red = titanic[['Age', 'Fare', 'Survived']]
```

```
[ ]: scaler = MinMaxScaler()
titanic_red[['Age', 'Fare']] = scaler.fit_transform(titanic_red[['Age', 'Fare']])
titanic_red.head()
```

/home/hakunamatata/.local/lib/python3.8/site-packages/pandas/core/frame.py:3678:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self[col] = igetitem(value, i)

```
[ ]:
```

	Age	Fare	Survived
0	0.273456	0.014151	0
1	0.473882	0.139136	1
2	0.323563	0.015469	1
3	0.436302	0.103644	1
4	0.436302	0.015713	0

```
[ ]:
```

```
x_trainr, x_testr, y_trainr, y_testr = train_test_split(titanic_red[['Age',
→ 'Fare']], titanic_red['Survived'], test_size = 0.2, stratify =
→ titanic_red['Survived'], random_state = 42)
print(x_trainr.shape, y_trainr.shape)
print(x_testr.shape, y_testr.shape)
```

```
(1045, 2) (1045,)
(262, 2) (262,)
```

```
[ ]: # https://scikit-learn.org/0.18/auto_examples/svm/plot_iris.html
xr, yr = x_trainr.values, y_trainr.values

h = .02 # step size in the mesh

# we create an instance of SVM and fit out data. We do not scale our
# data since we want to plot the support vectors
C = 1.0 # SVM regularization parameter
svc = SVC(kernel='linear')
rbf_svc = SVC(kernel='rbf')
poly_svc = SVC(kernel='poly')
sigmoid = SVC(kernel='sigmoid')

# create a mesh to plot in
x_min, x_max = xr[:, 0].min() - 0.1, xr[:, 0].max() + 0.1
y_min, y_max = xr[:, 1].min() - 0.1, xr[:, 1].max() + 0.1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))

# title for the plots
titles = ['SVC with linear kernel',
          'SVC with RBF kernel',
          'SVC with polynomial kernel',
          'SVC with sigmoid kernel']

for i, clf in enumerate((svc, rbf_svc, poly_svc, sigmoid)):
    # Plot the decision boundary. For that, we will assign a color to each
    # point in the mesh [x_min, x_max]x[y_min, y_max].
    plt.figure(figsize = (10,10))
    plt.subplot(2, 2, i + 1)
    plt.subplots_adjust(wspace=0.4, hspace=0.4)


    clf.fit(xr, yr)
    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])

    # Put the result into a color plot
    Z = Z.reshape(xx.shape)
```


```
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)

# Plot also the training points
plt.scatter(xr[:, 0], xr[:, 1], c=yr, cmap=plt.cm.coolwarm)
plt.xlabel('Age')
plt.ylabel('Fare')
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xticks(())
plt.yticks(())
plt.title(titles[i])

plt.show()
```




SVC with RBF kernel



SVC with polynomial kernel





```
[ ]: xr, yr = x_testr.values, y_testr.values

h = .02 # step size in the mesh

# create a mesh to plot in
x_min, x_max = -0.1, 1.1
y_min, y_max = -0.1, 1.1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))

# title for the plots
titles = ['SVC with linear kernel',
          'SVC with RBF kernel',
          'SVC with polynomial kernel',
          'SVC with sigmoid kernel']

for i, clf in enumerate((svc, rbf_svc, poly_svc, sigmoid)):
    # Plot the decision boundary. For that, we will assign a color to each
    # point in the mesh [x_min, x_max][y_min, y_max].
    plt.figure(figsize = (10,10))
    plt.subplot(2, 2, i + 1)
    plt.subplots_adjust(wspace=0.4, hspace=0.4)
```

```


Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])

# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)

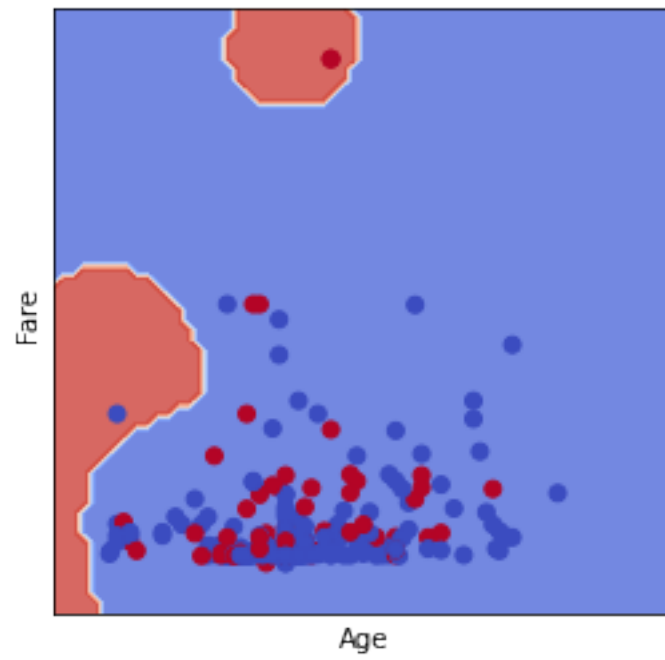
# Plot also the training points
plt.scatter(xr[:, 0], xr[:, 1], c=yr, cmap=plt.cm.coolwarm)
plt.xlabel('Age')
plt.ylabel('Fare')
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xticks(())
plt.yticks(())
plt.title(titles[i])

plt.show()


```

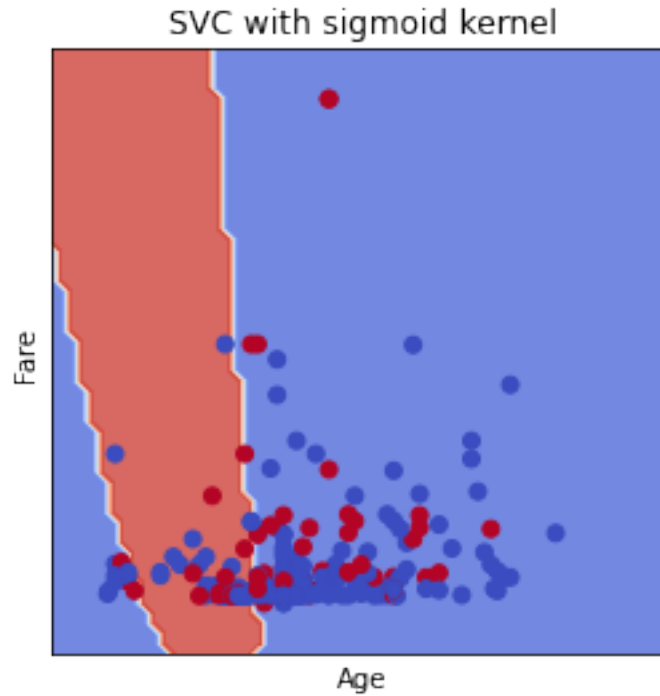


SVC with RBF kernel



SVC with polynomial kernel





```
[ ]: x_train = x_train.toarray()
     x_test = x_test.toarray()
```

```
[ ]: def log_kernel(xi, xj, d):
     s = 0
     for k in range(d):
         s += pow((xi[k] - xj[k]), d)
     return 0 if not s else -np.log(pow(s, 1/d)) + 1
```

```
[ ]: def logKernelGramMatrixFull(X1, X2):
     """(Pre)calculates Gram Matrix K"""

     log_matrix = np.zeros((X1.shape[0], X2.shape[0]))
     d = X1.shape[-1]
     for i, x1 in enumerate(X1):
         for j, x2 in enumerate(X2):
             x1 = x1.flatten()
             x2 = x2.flatten()
             log_matrix[i, j] = log_kernel(x1, x2, d)
     return log_matrix
```

```
[ ]: log_svc = SVC(C = 100, kernel = "precomputed")
     log_svc.fit(logKernelGramMatrixFull(x_train, x_train), y_train)
     y_pred = log_svc.predict(logKernelGramMatrixFull(x_test, x_train))
```

```
acc = accuracy_score(y_test, y_pred)
print(acc)
```

0.732824427480916

1.5 Task 3

```
[ ]: l,r = 1, 10000
i = 0
accs, cs = [],[]
while l < r:
    if i == 20: break
    mid = (r+l)/2
    lsvc = SVC(kernel = 'rbf', C = 1, gamma = 0.5)
    lsvc.fit(x_train, y_train)
    y_pred = lsvc.predict(x_test)
    lacc = accuracy_score(y_test, y_pred)
    rsvc = SVC(kernel = 'rbf', C = r, gamma = 0.5)
    rsvc.fit(x_train, y_train)
    y_pred = rsvc.predict(x_test)
    racc = accuracy_score(y_test, y_pred)
    print(l,r, lacc, racc)
    if lacc >= racc:
        r = round(mid,3)
        svc = lsvc
        cs.append(l)
    else:
        l = round(mid, 3)
        svc = rsvc
        cs.append(r)
    y_pred = svc.predict(x_test)
    accs.append(accuracy_score(y_test, y_pred))
    #cs.append(mid)

    i+=1
```

```
1 10000 0.7824427480916031 0.7595419847328244
1 5000.5 0.7824427480916031 0.7442748091603053
1 2500.75 0.7824427480916031 0.7404580152671756
1 1250.875 0.7824427480916031 0.7442748091603053
1 625.938 0.7824427480916031 0.7480916030534351
1 313.469 0.7824427480916031 0.7480916030534351
1 157.234 0.7824427480916031 0.7633587786259542
1 79.117 0.7824427480916031 0.7748091603053435
1 40.059 0.7824427480916031 0.7862595419847328
20.529 40.059 0.7786259541984732 0.7862595419847328
30.294 40.059 0.7824427480916031 0.7862595419847328
```


```

35.176 40.059 0.7824427480916031 0.7862595419847328
37.617 40.059 0.7824427480916031 0.7862595419847328
38.838 40.059 0.7862595419847328 0.7862595419847328
38.838 39.448 0.7862595419847328 0.7862595419847328
38.838 39.143 0.7862595419847328 0.7862595419847328
38.838 38.99 0.7862595419847328 0.7862595419847328
38.838 38.914 0.7862595419847328 0.7862595419847328
38.838 38.876 0.7862595419847328 0.7862595419847328
38.838 38.857 0.7862595419847328 0.7862595419847328

```

```
[ ]: plt.plot(accs)
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f6fd4b23310>]
```



1.6 Task 4

```

[ ]: l, r = 0.0001, 1
    i = 0
    accs, cs = [], []
    while l < r:
        if i == 20: break
        mid = (r+l)/2
        lsvc = SVC(kernel = 'rbf', C = 39, gamma = 1)
        lsvc.fit(x_train, y_train)
        y_pred = lsvc.predict(x_test)
        lacc = accuracy_score(y_test, y_pred)

```

```

rsvc = SVC(kernel = 'rbf', C = 39, gamma = r)
rsvc.fit(x_train, y_train)
y_pred = rsvc.predict(x_test)
racc = accuracy_score(y_test, y_pred)
print(l,r, lacc, racc)
if lacc >= racc:
    r = round(mid,3)
    svc = lsvc
    cs.append(l)
else:
    l = round(mid, 3)
    svc = rsvc
    cs.append(r)
y_pred = svc.predict(x_test)
accs.append(accuracy_score(y_test, y_pred))
#cs.append(mid)
i += 1

```


```

0.0001 1 0.7404580152671756 0.7519083969465649
0.5 1 0.7862595419847328 0.7519083969465649
0.5 0.75 0.7862595419847328 0.7824427480916031
0.5 0.625 0.7862595419847328 0.7824427480916031
0.5 0.562 0.7862595419847328 0.7824427480916031
0.5 0.531 0.7862595419847328 0.7862595419847328
0.5 0.516 0.7862595419847328 0.7862595419847328
0.5 0.508 0.7862595419847328 0.7862595419847328
0.5 0.504 0.7862595419847328 0.7862595419847328
0.5 0.502 0.7862595419847328 0.7862595419847328
0.5 0.501 0.7862595419847328 0.7862595419847328

```

```
[ ]: plt.plot(accs)
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f6fd4a99be0>]
```

1.7 Pre Processing on Forest Cover Dataset

```
[ ]: forest = pd.read_csv("covtype.csv")
forest.head()
```

```
[ ]:
Elevation  Aspect  Slope  Horizontal_Distance_To_Hydrology  \
0         2596     51      3                               258
1         2590     56      2                               212
2         2804    139      9                               268
3         2785    155     18                               242
4         2595     45      2                               153

Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \
0                               0                               510
1                              -6                               390
2                               65                              3180
3                              118                              3090
4                              -1                               391

Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \
0             221             232             148
1             220             235             151
2             234             238             135
3             238             238             122
4             220             234             150
```

	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	Soil_Type33	\
0	6279	...	0	0	
1	6225	...	0	0	
2	6121	...	0	0	
3	6211	...	0	0	
4	6172	...	0	0	

	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	Soil_Type38	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Soil_Type39	Soil_Type40	Cover_Type
0	0	0	5
1	0	0	5
2	0	0	2
3	0	0	2
4	0	0	5

[5 rows x 55 columns]

```
[ ]: forest.describe()
```

```
[ ]:
      Elevation      Aspect      Slope  \
count  581012.000000  581012.000000  581012.000000
mean    2959.365301    155.656807    14.103704
std      279.984734    111.913721     7.488242
min     1859.000000     0.000000     0.000000
25%     2809.000000     58.000000     9.000000
50%     2996.000000    127.000000    13.000000
75%     3163.000000    260.000000    18.000000
max     3858.000000    360.000000    66.000000
```

	Horizontal_Distance_To_Hydrology	Vertical_Distance_To_Hydrology	\
count	581012.000000	581012.000000	
mean	269.428217	46.418855	
std	212.549356	58.295232	
min	0.000000	-173.000000	
25%	108.000000	7.000000	
50%	218.000000	30.000000	
75%	384.000000	69.000000	
max	1397.000000	601.000000	

	Horizontal_Distance_To_Roadways	Hillshade_9am	Hillshade_Noon	\
--	---------------------------------	---------------	----------------	---

count	581012.000000	581012.000000	581012.000000
mean	2350.146611	212.146049	223.318716
std	1559.254870	26.769889	19.768697
min	0.000000	0.000000	0.000000
25%	1106.000000	198.000000	213.000000
50%	1997.000000	218.000000	226.000000
75%	3328.000000	231.000000	237.000000
max	7117.000000	254.000000	254.000000

	Hillshade_3pm	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	\
count	581012.000000		581012.000000	...	581012.000000
mean	142.528263		1980.291226	...	0.090392
std	38.274529		1324.195210	...	0.286743
min	0.000000		0.000000	...	0.000000
25%	119.000000		1024.000000	...	0.000000
50%	143.000000		1710.000000	...	0.000000
75%	168.000000		2550.000000	...	0.000000
max	254.000000		7173.000000	...	1.000000

	Soil_Type33	Soil_Type34	Soil_Type35	Soil_Type36	\
count	581012.000000	581012.000000	581012.000000	581012.000000	
mean	0.077716	0.002773	0.003255	0.000205	
std	0.267725	0.052584	0.056957	0.014310	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	Soil_Type37	Soil_Type38	Soil_Type39	Soil_Type40	\
count	581012.000000	581012.000000	581012.000000	581012.000000	
mean	0.000513	0.026803	0.023762	0.015060	
std	0.022641	0.161508	0.152307	0.121791	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	Cover_Type
count	581012.000000
mean	2.051471
std	1.396504
min	1.000000
25%	1.000000
50%	2.000000
75%	2.000000

max 7.000000

[8 rows x 55 columns]

```
[ ]: X1 = forest.copy()
X1[['Elevation', 'Aspect', 'Slope', 'Horizontal_Distance_To_Hydrology',
     'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
     'Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm',
     'Horizontal_Distance_To_Fire_Points', 'Wilderness_Area1',
     'Wilderness_Area2', 'Wilderness_Area3', 'Wilderness_Area4']] = scaler.
↳fit_transform(X1[['Elevation', 'Aspect', 'Slope',
↳'Horizontal_Distance_To_Hydrology',
     'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
     'Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm',
     'Horizontal_Distance_To_Fire_Points', 'Wilderness_Area1',
     'Wilderness_Area2', 'Wilderness_Area3', 'Wilderness_Area4']])
X1 = X1.drop(columns = ['Cover_Type'])
Y1 = forest['Cover_Type']
print(Y1.shape)
X1.head()
```

(581012,)

```
[ ]:  Elevation    Aspect    Slope  Horizontal_Distance_To_Hydrology  \
0    0.368684  0.141667  0.045455                                0.184681
1    0.365683  0.155556  0.030303                                0.151754
2    0.472736  0.386111  0.136364                                0.191840
3    0.463232  0.430556  0.272727                                0.173228
4    0.368184  0.125000  0.030303                                0.109520

      Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \
0                                0.223514                        0.071659
1                                0.215762                        0.054798
2                                0.307494                        0.446817
3                                0.375969                        0.434172
4                                0.222222                        0.054939

      Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \
0          0.870079          0.913386          0.582677
1          0.866142          0.925197          0.594488
2          0.921260          0.937008          0.531496
3          0.937008          0.937008          0.480315
4          0.866142          0.921260          0.590551

      Horizontal_Distance_To_Fire_Points  ...  Soil_Type31  Soil_Type32  \
0                                0.875366  ...           0           0
1                                0.867838  ...           0           0
```

2	0.853339	...	0	0
3	0.865886	...	0	0
4	0.860449	...	0	0

	Soil_Type33	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Soil_Type38	Soil_Type39	Soil_Type40
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 54 columns]

```
[ ]: _, subset_x, _, subset_y = train_test_split(X1, Y1, stratify = Y1, test_size = 0.2, random_state = 42)
x_train1, x_test1, y_train1, y_test1 = train_test_split(subset_x, subset_y, stratify = subset_y, test_size = 0.1, random_state = 42)
print(x_train1.shape, x_test1.shape)
print(y_train1.shape, y_test1.shape)
```

```
(104582, 54) (11621, 54)
(104582,) (11621,)
```

1.8 Task 1

```
[ ]: svc = SVC(kernel='linear')
ovr = OneVsRestClassifier(svc)
ovr.fit(x_train1, y_train1)
y_pred1 = ovr.predict(x_test1)
acc1 = accuracy_score(y_test1, y_pred1)
```


```
[ ]: cm = confusion_matrix(y_test1, y_pred1, labels=ovr.classes_)
print(cm)
```

```
[[3000 1154    4    0    2    2   75]
 [1038 4462  125    0    1   23   18]
 [    0    68  635    2    0   10    0]
 [    0     0   44   10    0    1    0]
 [    3   168    9    0    0   10    0]
 [    0   133  197    0    0   17    0]
```

```
[ 206    0    1    0    0    0  203]]
```

```
[ ]: ax = sns.heatmap(cm, annot=True, fmt='g')
ax.set_title('Confusion Matrix')
ax.set_xlabel('Predicted Labels')
ax.set_ylabel('Actual Labels')
ax.xaxis.set_ticklabels(ovr.classes_)
ax.yaxis.set_ticklabels(ovr.classes_)
```

```
[ ]: [Text(0, 0.5, '1'),
      Text(0, 1.5, '2'),
      Text(0, 2.5, '3'),
      Text(0, 3.5, '4'),
      Text(0, 4.5, '5'),
      Text(0, 5.5, '6'),
      Text(0, 6.5, '7')]
```




```
[ ]: ovo = OneVsOneClassifier(svc)
ovo.fit(x_train1,y_train1)
y_pred2 = ovo.predict(x_test1)
acc2 = accuracy_score(y_test1,y_pred2)
```

```
[ ]: cm = confusion_matrix(y_test1,y_pred2,labels=ovo.classes_)
print(cm)
```

```
[[3040 1122    1    0    0    2   72]
 [1056 4498   90    0    0   11   12]
 [    0    67  628    5    0   15    0]
 [    0    0   43   11    0    1    0]
 [    0   186    4    0    0    0    0]
 [    0    87  206    3    0   51    0]
 [ 188    1    0    0    0    0  221]]
```

```
[ ]: ax = sns.heatmap(cm, annot=True, fmt='g')
ax.set_title('Confusion Matrix')
ax.set_xlabel('Predicted Labels')
ax.set_ylabel('Actual Labels')
ax.xaxis.set_ticklabels(ovo.classes_)
ax.yaxis.set_ticklabels(ovo.classes_)
```

```
[ ]: [Text(0, 0.5, '1'),
      Text(0, 1.5, '2'),
      Text(0, 2.5, '3'),
      Text(0, 3.5, '4'),
      Text(0, 4.5, '5'),
      Text(0, 5.5, '6'),
      Text(0, 6.5, '7')]
```



```
[ ]: print("OvR Acc Score: ",acc1)
      print("OvO Acc Score: ",acc2)
```

OvR Acc Score: 0.7165476292917994

OvO Acc Score: 0.7270458652439549